

### Problem Set 3

CS265, Autumn 2022-23

Due: October 21 (Friday)

Group members: Rylan Schaeffer, Various people at office hours (mainly the CAs)

Please follow the homework policies on the course website.

---

#### 1. (11 pt.) Aggregating Guesses

In this problem, we'll consider several different settings where we are aggregating a large number of noisy, unbiased estimates. Suppose a class has  $n$  students. Each student is asked to estimate the current temperature. Assume that they each provide independent, unbiased estimates, with  $X_i$  denoting the  $i$ th student's guess. Let  $v_i$  denote  $\text{Var}[X_i]$ .

- (a) (4 pt.) Suppose we know each of the  $v_i$ 's and decide to compute a weighted combination  $Z = \sum_i w_i X_i$ , where the weights  $w_i \geq 0$  are chosen so as to minimize the variance of  $Z$ , subject to  $\sum_i w_i = 1$ . What are those optimal weights as a function of the  $v_i$ 's, and roughly how accurate will  $Z$  be? Please give an answer of the form: "with probability at least 0.9,  $Z$  will be within *blah* of the true temperature, where *blah* is a function of the  $v_i$ 's."

- (b) (5 pt.) For this part, assume each  $X_i$  is drawn from a normal (Gaussian) distribution, whose mean is the true temperature, and whose variance is 1. Roughly how accurate should we expect the *median* of the  $n$  guesses to be? As above, please give an answer of the form: "with probability at least 0.9, the median of the  $X_i$ 's will be within *blah* of the true temperature," where *blah* is a function of  $n$ . Your value of *blah* should be accurate up to a constant factor and use big-Oh notation, for example  $O(1/n^{3/4})$  or something like that.

[HINT: The following basic fact about a Gaussian should be helpful, and is the only property of a Gaussian that you will need: if  $Y$  is a Gaussian with mean  $\mu$  and variance 1, for any  $\epsilon \in (0, 1/2)$   $\Pr[Y < \mu - \epsilon] = \Pr[Y > \mu + \epsilon] < 1/2 - 0.3\epsilon$ .]

- (c) (2 pt.) Answer the same question as above for the *mean* of the  $n$  values  $X_i$ . How do your answers compare?
- (d) (0 pt.) [Optional: This is a research-level problem.] As above, suppose each  $X_i$  is independently drawn from a normal distribution whose mean is the true temperature, and variance  $v_i$ . Assume you know the (multi)set of the  $v_i$ 's, but you don't know which variance corresponds to which guess. How well should you expect to do, and is there an efficient algorithm that achieves this?

- (e) (0 pt.) [Optional: This is a research-level problem.] Suppose we are in the setting above, but don't know anything about the variances. What is a near-optimal algorithm, and how well will it do, as a function of the (unknown) list of variances  $v_1, \dots$ ?

[HINT: Note that if two  $X_i$ 's are identical (or super, super close) then we know that two of the variances are 0 (or really, really small), and hence either of those  $X_i$ 's would give an extremely accurate guess, no matter what the other  $n - 2$  guesses are...]

**SOLUTION:**

(a) First, we simplify the term of interest:

$$\mathbb{V}[Z] = \mathbb{V}[\sum_i w_i X_i] = \sum_i w_i^2 \mathbb{V}[X_i] = \sum_i w_i^2 v_i$$

Then we define the Lagrangian:

$$\mathcal{L}(w_i, \dots, w_n, \lambda) := \sum_i w_i^2 v_i + \lambda(1 - \sum_i w_i)$$

Computing the first-order stationary points yields:

$$\begin{aligned} \partial_{w_j} \mathcal{L} = 0 &= 2w_j v_j - \lambda \Rightarrow w_j = \frac{\lambda}{2v_j} \\ \partial_{\lambda} \mathcal{L} = 0 &= 1 - \sum_i w_i \Rightarrow 1 = \sum_i w_i \end{aligned}$$

Rearranging yields  $\lambda = 2/\sum_i v_i$ , and plugging back into  $w_j$  yields  $w_j = 1/(v_j \sum_i \frac{1}{v_i})$ . How accurate will  $Z$  be? Let's first compute the variance in  $Z$ , then apply Chebyshev's:

$$\mathbb{V}[Z] = \mathbb{V}[\sum_i w_i X_i] = \frac{1}{\sum_i \frac{1}{v_i}}$$

Then, per Chebyshev's inequality, we know that:

$$\mathbb{P}[|Z - \mu^*| \geq k\sigma] \leq \frac{1}{k^2} = 0.1$$

So with probability at least 0.9,  $Z$  will be within  $\sqrt{\frac{10}{\sum_i \frac{1}{v_i}}}$  of the true mean.

(b) Let  $m := \text{median}(X_1, \dots, X_n)$ ,  $\epsilon \geq 0$  denote our (unknown) “blah” function of  $n$  and  $\mu^*$  denote the true (unknown) temperature. We want the probability that  $m \in [\mu^* - \epsilon, \mu^* + \epsilon]$  to be at least 0.9. In this problem, we'll find it easier to work with 1 minus the complement:

$$0.9 \leq \mathbb{P}[m \in \mu^* \pm \epsilon] = 1 - \mathbb{P}[m \leq \mu^* - \epsilon] - \mathbb{P}[m \geq \mu^* + \epsilon]$$

By symmetry,  $\mathbb{P}[m \leq \mu^* - \epsilon] = \mathbb{P}[m \geq \mu^* + \epsilon]$ , so we can work with:

$$0.9 \leq \mathbb{P}[m \in \mu^* \pm \epsilon] = 1 - 2\mathbb{P}[m \geq \mu^* + \epsilon]$$

Rearranging, we want to find  $\epsilon$  such that the following holds:

$$\mathbb{P}[m \geq \mu^* + \epsilon] \leq 0.05$$

The median will be greater than the mean by  $\epsilon$  if at least  $n/2$  elements are greater than the mean by  $\epsilon$ . Define the indicator variable  $I_i := X_i > \mu^* + \epsilon$  and the sum of the indicator variables  $S := \sum_i I_i$ . Our (sub)goal is now to choose  $\epsilon \in (0, \frac{1}{2})$  such that  $\mathbb{P}[S \geq \frac{n}{2}] \leq 0.05$ . Under Chernoff's Inequality, we have:

$$\mathbb{P}[S \geq n/2] \leq \exp(-\delta^2 n(1/2 - 0.3\epsilon)/3)$$

where  $\delta > \frac{1}{1-0.6\epsilon} - 1 = \frac{0.6\epsilon}{1-0.6\epsilon} > 0.6\epsilon$ . Simplifying:

$$\mathbb{P}[S \geq n/2] \leq \exp(-\frac{6}{100}n\epsilon^2)$$

We want the RHS to be 0.05. Choosing  $\epsilon = O(n^{-1/2})$  satisfies this. In conclusion, with probability  $\geq 0.9$ , the median will be within  $O(1/\sqrt{n})$  of the true mean  $\mu^*$ .

- (c) Again, let  $\mu^*$  denote the true temperature and define  $\bar{X} := \frac{1}{n} \sum_i X_i$ . Because a linear combination of Gaussians is itself Gaussian, all we need to calculate are the first two moments. The mean is:

$$\mathbb{E}[\bar{X}] = \frac{1}{n} \sum_i \mathbb{E}[X_i] = n \frac{1}{n} \mu^* = \mu^*$$

The variance is:

$$\mathbb{V}[\bar{X}] = \frac{1}{n^2} \sum_i \mathbb{V}[X_i] = \frac{1}{n^2} n = \frac{1}{n}$$

Thus,  $\bar{X} \sim \mathcal{N}(\mu^*, \frac{1}{n})$ . This means that with probability 0.9,  $\bar{X}$  is within the interval  $\mu^* \pm 1.65/\sqrt{n}$  of the true mean, meaning that with probability 0.9,  $\bar{X}$  will be within  $1.65/\sqrt{n}$  of the true temperature.

## 2. (11 pt.) Concentration without Independence

A computer system has  $n$  different failure modes, each of which happens with a small probability. Fortunately, the system is designed to be sufficiently robust in the following sense: as long as less than half of the failures occur, things are fine; otherwise, a large-scale crash will happen. We want to make sure that the probability of crashing is small enough.

To model the above scenario, we define  $n$  Bernoulli random variables  $X_1, \dots, X_n$ . Each  $X_i$  is the indicator of the  $i$ -th failure mode, i.e.,  $X_i = 1$  if failure  $i$  occurs and  $X_i = 0$  otherwise. Our goal is to upper bound the probability  $\Pr[\sum_{i=1}^n X_i \geq n/2]$ .

- (a) **(2 pt.)** Let's first assume that the  $n$  failure events are independent and the probability of each failure is at most  $1/3$ . Formally, we have:

**Assumption 1.**  $\Pr[X_i = 1] \leq 1/3$  for every  $i \in [n]$  and  $X_1, \dots, X_n$  are independent.

Prove that under Assumption 1, for some constant  $C > 0$  that does not depend on  $n$ ,

$$\Pr\left[\sum_{i=1}^n X_i \geq n/2\right] \leq e^{-Cn}. \quad (1)$$

Thus, the probability of a crash is exponentially small in  $n$ .

[**HINT:** Feel free to use (without proof) any of the Chernoff bounds in lecture note #5 (including Theorem 2 and Corollaries 5 and 6) and also the inequality  $\frac{e^\delta}{(1+\delta)^{1+\delta}} \leq e^{-\delta^2/3}$  for  $\delta \in [0, 1]$ . ]

- (b) **(1 pt.)** Now we decide that Assumption 1 is too unrealistic, since many of the failure modes are known to be strongly correlated. Show that only assuming  $\Pr[X_i = 1] \leq 1/3$  (and not the independence), the probability of crashing can be as large as  $\Omega(1)$ . In particular, prove that for any  $n \geq 1$ , there exist random variables  $X_1, \dots, X_n$  that satisfy: (1)  $\Pr[X_i = 1] \leq 1/3$  for every  $i \in [n]$ ; (2)  $\Pr[\sum_{i=1}^n X_i \geq n/2] \geq 1/3$ .
- (c) **(2 pt.)** Let's try the following relaxation of Assumption 1, which states that the probability for  $k$  different failures to occur simultaneously is exponentially small in  $k$ :

**Assumption 2.** For any  $S \subseteq [n]$ ,  $\Pr[X_i = 1 \text{ for all } i \in S] \leq (1/3)^{|S|}$ .

Show that Assumption 2 is strictly weaker than Assumption 1 by proving: (1) Assumption 1 implies Assumption 2; (2) the implication on the other direction does not hold, i.e., there exist some  $n$  and  $X_1, \dots, X_n$  that satisfy Assumption 2 but not Assumption 1.

[**HINT:** For (2), there exists a counterexample for  $n = 2$ . ]

- (d) **(6 pt.)** Prove that under Assumption 2, inequality (1) holds for some constant  $C > 0$ . In your proof, you can appeal to the proof of the Chernoff bounds from lecture videos/notes if you need to write it out verbatim at some point. For example, if you manage to upper bound  $\Pr[\sum_{i=1}^n X_i \geq n/2]$  by an expression involving the moment-generating function of some random variable  $Y$  that is the sum of  $n$  independent Bernoulli random variables, you can simply say that “the rest of the proof is exactly the proof of Theorem 2 from Lecture #5”.

[**HINT:** Consider independent Bernoulli random variables  $Y_1, \dots, Y_n$  with  $\Pr[Y_i = 1] = 1/3$  for each  $i \in [n]$ . For distinct indices  $i, j, \ell \in [n]$ , does  $\mathbb{E}[X_i X_j X_\ell] \leq \mathbb{E}[Y_i Y_j Y_\ell]$  hold? Can you extend your proof of the inequality to the case with repeating indices? ]

[**HINT:** Let  $X = \sum_{i=1}^n X_i$  and  $Y = \sum_{i=1}^n Y_i$ . What can we say about  $\mathbb{E}[X^k]$  and  $\mathbb{E}[Y^k]$  for integer  $k \geq 0$ ? Considering the identity  $e^z = \sum_{k=0}^{+\infty} \frac{z^k}{k!}$ , what can we say about  $\mathbb{E}[e^{tX}]$  and  $\mathbb{E}[e^{tY}]$  for any  $t > 0$ ? ]

- (e) **(0 pt.) [Optional: this won't be graded.]** Can you construct counterexamples for Part 2b that satisfy *pairwise independence* but have a crashing probability of  $\Omega(1/n)$ ? Formally, prove that there exists  $C > 0$  such that for any  $n \geq 2$ , there exist  $X_1, \dots, X_n$  that satisfy: (1)  $\Pr[X_i = 1] \leq 1/3$ ; (2)  $X_i$  and  $X_j$  are independent for distinct  $i, j \in [n]$ ; (3)  $\Pr[\sum_{i=1}^n X_i \geq n/2] \geq C/n$ .

[**NOTE:** This shows that unlike Chebyshev's inequality, Chernoff bounds do not hold if we only assume pairwise independence. ]

[**HINT:** Recall pairwise independent hash functions if you have seen them before. You can use the Bertrand-Chebyshev theorem, which states that for any integer  $n \geq 1$ , there exists a prime number  $p$  with  $n < p < 2n$ . ]

**SOLUTION:**

(a) Let  $X^* \sim \text{Bern}(1/3)$ . Then under Assumption 1:

$$\begin{aligned}\mathbb{P}[\sum_i X_i \geq \frac{n}{2}] &\leq \mathbb{P}[\sum_i X^* \geq \frac{n}{2}] \\ &= \mathbb{P}[\sum_i X^* \geq (1 + \frac{1}{2})\frac{n}{3}]\end{aligned}$$

Applying Corollary 5 from Lecture Note #5, we have:

$$\begin{aligned}\mathbb{P}[\sum_i X_i \geq \frac{n}{2}] &\leq e^{-\mu\delta^2/3} \\ &= e^{-(n/3)(1/2)^2/3} \\ &= e^{-n/36}\end{aligned}$$

(b) Suppose  $\{X_i\}$  are all perfectly correlated and  $X_i \sim \text{Bern}(1/3)$ . Then

$$\mathbb{P}[\sum_i X_i \geq \frac{n}{2}] = \frac{1}{3}$$

(c) Suppose Assumption 1 is true. Let  $S \subseteq [n]$  with cardinality  $1 \leq |S| \leq n$ . Then:

$$\mathbb{P}[X_{s_1} = 1, X_{s_2} = 1, \dots, X_{s_{|S|}} = 1] = \prod_{s=1}^{|S|} \mathbb{P}[X_s = 1] \leq \prod_{s=1}^{|S|} \left(\frac{1}{3}\right) = \left(\frac{1}{3}\right)^{|S|}$$

Therefore Assumption 1  $\Rightarrow$  Assumption 2. However, the reverse does not hold. Let  $n = 2$  and suppose Assumption 2 is true. Suppose  $X_1 \sim \text{Bern}(\frac{1}{3})$  and  $X_2 \sim X_1 * \text{Bern}(\frac{1}{3})$ ; in other words,  $X_2$  is constructed by flipping a coin with probability  $1/3$ , but only keeping 1 if  $X_1$  is also 1. Note that  $P(X_1 = 1) = 1/3$ ,  $P(X_2 = 1) < 1/3$  and  $P(X_1 = 1, X_2 = 1) \leq (1/3)^2$ , thus satisfying Assumption 2, but breaking the independence assumption because  $P(X_1 = 1, X_2 = 1) \neq P(X_1 = 1)P(X_2 = 1)$ .

(d) Define  $X := \sum_i X_i$  and  $Y := \sum_i Y_i$ . Let's first prove the two hints. Hint 1 says that for  $i \neq j \neq \ell$ ,  $\mathbb{E}[X_i X_j X_\ell] \leq \mathbb{E}[Y_i Y_j Y_\ell]$ . This is true because  $\mathbb{E}[X_i X_j X_\ell] = \mathbb{P}[X_i = 1, X_j = 1, X_\ell = 1] \leq (1/3)^3$  under Assumption 2, whereas  $\mathbb{E}[Y_i Y_j Y_\ell] = (1/3)^3$ . The same logic holds morally for repeat indices (with slight modification to avoid double-counting the exponent because e.g.  $P(X_i = 1, X_i = 1) = P(X_i = 1)$ ). This then implies that  $\forall k \geq 0, \mathbb{E}[X^k] \leq \mathbb{E}[Y^k]$ , and can be used to prove Hint 2:

$$\begin{aligned}\mathbb{E}[e^{tX}] &= \mathbb{E}[1 + \frac{tX}{1} + \frac{(tX)^2}{2} + \dots] \\ &= 1 + \frac{\mathbb{E}[(tX)]}{1} + \frac{\mathbb{E}[(tX)^2]}{2} + \dots \\ &\leq 1 + \frac{\mathbb{E}[tY]}{1} + \frac{\mathbb{E}[(tY)^2]}{2} + \dots \\ &= \mathbb{E}[e^{tY}]\end{aligned}$$

We now return to our goal:

$$\mathbb{P}[X \geq \frac{n}{2}] = \mathbb{P}[e^{tX} \geq e^{tn/2}] \leq \frac{\mathbb{E}[e^{tX}]}{e^{tn/2}} \leq \frac{\mathbb{E}[e^{tY}]}{e^{tn/2}}$$

The rest of the proof is exactly the proof of Theorem 2 from Lecture # 5.

### 3. (8 pt.) Processes and CPUs

Suppose that in a distributed system, we have  $N$  CPUs and  $P$  processes. Each process is independently and uniformly allocated to a CPU. However, if multiple processes are allocated to the same CPU, the CPU will choose one of them at random to complete; the remaining processes allocated to that CPU will not be completed.

- (a) (2 pt.) What is expected number of processes that will be completed?
- (b) (6 pt.) Suppose that  $P \geq N$ , and denote the total number of completed processes by  $C$ . Use **Poissonization** to prove that  $\Pr[C \leq \frac{N}{2}] \leq e^{-\Omega(N)}$ .  
[Note: There may be ways to do this problem that don't involve Poissonization, but we want you to use it to get practice with it. That is, you should prove the statement by analyzing the case when the number of processes is an appropriate Poisson random variable. Don't forget the de-Poissonization step!]
- (c) (0 pt.) [Optional: this won't be graded.] Let  $\mu$  be your answer from part (a). Under what conditions on  $P$  and  $N$  can you use Poissonization to show that  $\Pr[C \leq (1 - \delta)\mu] \leq e^{-\Omega_\delta(N)}$  for any  $\delta > 0$ , where the  $\Omega_\delta$  notation means that you are allowed to have constants that depend on  $\delta$  hidden inside the big-Omega.

#### SOLUTION:

- i. The number of completed processes is given by:

$$\begin{aligned} \# \text{ completed processes} &= \sum_i \mathbb{I}[\text{at least 1 process assigned to } i\text{th CPU}] \\ &= \sum_i 1 - \sum_i \mathbb{I}[0 \text{ processes assigned to } i\text{th CPU}] \\ &= n - n\mathbb{P}[0 \text{ processes assigned to } i\text{th CPU}] \\ &= n \left( 1 - \left( \frac{n-1}{n} \right)^P \right) \end{aligned}$$

- ii. Let's start with the Poissonization step. Suppose we draw  $k \sim \text{Pois}(2P)$  processes, then the number of processes per CPU is:

$$X_n \sim \text{Pois}(2P)$$

Define  $C$  as the number of complete processes. Then

$$C = \sum_{n=1}^N \mathbb{I}[X_n \geq 1] = \sum_n (1 - \mathbb{I}[X_n = 0]) = N - \sum_n \mathbb{I}[X_n = 0]$$

In this Poissonified setting, with independence, we can bound the probability  $C \leq N/2$  using Chernoff:

$$\mathbb{P}[C \leq N/2] = \mathbb{P}[C \leq (1 - \delta)\mathbb{E}[C]] \leq \exp(-\delta^2\mathbb{E}[C]/2)$$

where  $\delta := 1 - N/2\mathbb{E}[C]$ . The expected number of completed processes is:

$$\mathbb{E}[C] = N - N\mathbb{P}[Pois(2P/N) = 0] = N(1 - \exp(-2P/N))$$

Simplifying :

$$\begin{aligned} \exp(-\delta^2\mathbb{E}[C]/2) &= \exp\left(-\left(1 - \frac{N}{2\mathbb{E}[C]}\right)^2\mathbb{E}[C]/2\right) \\ &= \exp\left(-\frac{1}{2}\left(\mathbb{E}[C] - N + \frac{N^2}{4\mathbb{E}[C]}\right)\right) \\ &= \exp(-\Omega(N)) \end{aligned}$$

where the last statement is true because (with  $P \geq N$ ):  $\mathbb{E}[C] < N$ ,  $N = N$  and  $N^2/4\mathbb{E}[C] \leq N$ .

Next, let's de-Poissonify. Suppose that instead of drawing  $k \sim Pois(2P)$  processes, consider if we had exactly  $P$  processes. There are two possible cases. In the first case of  $k \geq P$ , then  $\mathbb{P}[C \leq \frac{N}{2}]$  must shrink because drawing more processes than  $P$  means that the probability of empty CPUs falls, and thus the same bound holds. In the second case of  $k < P$ , we can use Fact 3 from Lecture #6 to show that this is exceptionally unlikely:

$$\mathbb{P}[|k - 2P| \geq P] \leq 2\exp\left(-\frac{P^2}{2(P + 2P)}\right) \leq 2\exp(-P/6) = o(1)$$